

# Notebook for the Project

## Introduction

Do MNIST-1D, compare CNN, RandomForest with PCA, LinearRegression with PCA

Goal: Compare the relation of accuracy to training and prediction time. Expected outcome: CNN is the most accurate, but will be a lot slower than the others.

## Dependencies

```
In [1]: # !pip install torch
# !pip install pandas
# !pip install scikit-learn
# !pip install tensorboard
# !pip install seaborn
# !pip install torchvision
```

## Imports

```
In [2]: import torch
from torch import autograd
from torch import nn
from torch.utils.data import TensorDataset, DataLoader
from torchvision.datasets import MNIST
from torchvision import transforms

from matplotlib import pyplot as plt
import numpy as np
import pandas as pd
import seaborn as sns
import numpy as np
from tqdm import tqdm
from copy import deepcopy

from collections import OrderedDict
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from torch.utils.tensorboard import SummaryWriter
import sklearn.metrics as metrics
from pandas import Series
from typing import Union
import json
from sklearn.metrics import confusion_matrix, accuracy_score

from sklearn.decomposition import PCA
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import RandomizedSearchCV
```

```
/Users/veikko/Documents/Kouluhommia/MachineLearning/aalto-machine-learning/.venv/lib/pyth
hon3.9/site-packages/tqdm/auto.py:22: TqdmWarning: IPProgress not found. Please update ju
pyter and ipywidgets. See https://ipywidgets.readthedocs.io/en/stable/user_install.html
from .autonotebook import tqdm as notebook_tqdm
```

## Helper functions

```
In [3]: def train(train_loader, model, optimizer, criterion, device, detect_bad_gradients=False,
        """
        Trains PyTorch model for one epoch in batches.

        Args:
            train_loader: Data loader for training set.
            model: Neural network model.
            optimizer: Optimizer (e.g. SGD).
            criterion: Loss function (e.g. cross-entropy loss).
        """

        avg_loss = 0

        model.train()

        # Iterate through batches
        for i, data in enumerate(train_loader):
            # Get the inputs; data is a list of [inputs, labels]
            inputs, labels = data

            # Move data to target device
            inputs, labels = inputs.to(device), labels.to(device)

            # Zero the parameter gradients
            optimizer.zero_grad()

            # Forward + backward + optimize
            outputs = model(inputs)
            loss = criterion(outputs, labels) # Compute RMSE from MSE
            if detect_bad_gradients:
                with autograd.detect_anomaly():
                    loss.backward()
            else:
                loss.backward()
            if clip_grad_norm:
                grad_norm = torch.nn.utils.clip_grad_norm_(model.parameters(), clip_grad_norm)
            optimizer.step()

            # Keep track of loss (MSE) and r2
            avg_loss += torch.sqrt(loss)

        return avg_loss / len(train_loader)

def test(test_loader, model, criterion, device):
    """
    Evaluates network in batches.

    Args:
        test_loader: Data loader for test set.
        model: Neural network model.
        criterion: Loss function (e.g. cross-entropy loss).
    """
```

```

avg_loss = 0

model.eval()

# Use torch.no_grad to skip gradient calculation, not needed for evaluation
with torch.no_grad():

    # Iterate through batches
    all_predictions = []
    all_labels = []

    for data in test_loader:
        # Get the inputs; data is a list of [inputs, labels]
        inputs, labels = data

        # Move data to target device
        inputs, labels = inputs.to(device), labels.to(device)

        # Forward pass
        outputs = model(inputs)
        loss = torch.sqrt(criterion(outputs, labels)) # Compute RMSE from MSE

        all_predictions.extend(outputs.detach().numpy())
        all_labels.extend(labels.detach().numpy())

        # Keep track of loss (MSE) and r2
        avg_loss += loss

    return avg_loss / len(test_loader)
    # Track the average loss and the r2 of the last batch

def run_torch(model, train_set, val_set, test_set, log_comment="", log_hparams=False, writer=None):
    """
    Run a test
    """
    try:
        if writer is None:
            # Create a writer to write to Tensorboard
            writer = SummaryWriter(comment=log_comment)
            writer.add_text("run_params", json.dumps(config, indent=2))

        # Create the dataloaders
        train_loader = DataLoader(
            train_set, batch_size=config["batch_size"], shuffle=False, num_workers=0, pin_memory=True
        )
        val_loader = DataLoader(
            val_set, batch_size=config["batch_size"], shuffle=False, num_workers=0, pin_memory=True
        )
        test_loader = DataLoader(
            test_set, batch_size=config["batch_size"], shuffle=False, num_workers=0, pin_memory=True
        )

        # Create loss function and optimizer
        if config["loss"] == "MSE" or config["loss"] == "RMSE":
            criterion = nn.MSELoss()
        else:
            raise ValueError(f"Loss {config['loss']} not recognized.")
    except KeyboardInterrupt:
        print("KeyboardInterrupt")

```

```

optimizer = torch.optim.Adam(
    model.parameters(), lr=config["optimizer"]["lr"], weight_decay=config["optim
")
if config["lr_scheduler"]:
    scheduler = torch.optim.lr_scheduler.MultiStepLR(optimizer=optimizer, **conf

# Use GPU if available
device = "cuda" if torch.cuda.is_available() else "cpu"

model = model.to(device)

patience = config.get("early_stopping_patience", torch.inf)
best_model = None
best_loss = np.inf
counter = 0

print("Starting initial training")
for epoch in tqdm(range(config["epochs"])):
    # Train on data
    train_loss = train(
        train_loader, model, optimizer, criterion, device, config["detect_bad_gr
    )
    # After training set eval mode on
    model.eval()
    # Test on data
    val_loss = test(val_loader, model, criterion, device)
    test_loss = test(test_loader, model, criterion, device)

    if config["lr_scheduler"]:
        scheduler.step()

    # Write metrics to Tensorboard
    writer.add_scalars("Loss", {"Train_loss": train_loss, "Val_loss": val_loss,

    if log_hparams:
        report_metrics = {
            "hparam/test_loss": test_loss,
            "hparam/train_loss": train_loss,
        }
        writer.add_hparams(log_hparams, report_metrics, run_name=log_comment)

    # Early stopping
    if best_loss > val_loss.detach().numpy():
        best_loss = val_loss.detach().numpy()
        counter = 0
        best_model = deepcopy(model)
    else:
        counter += 1
        if counter > patience:
            print("Initiating early stopping")
            if best_model is not None:
                print("Restoring best weights")
                model = best_model
            break

print("\nTraining Finished.")
writer.flush()
writer.close()

# Finally, use the model to predict the train, validation and test sets
except KeyboardInterrupt:

```

```

        print("Interrupted")
    print("Gathering final predictions")

    if not log_hparams:
        results, predictions, model = gather_results(model, train_loader, val_loader, te

        return results, predictions, model
    else:
        return

def gather_results(model, train_loader, val_loader, test_loader):
    """
    Gather the results for train, val and test sets.
    Returns:
        results, predictions, model
    """
    model.eval()
    with torch.no_grad():

        y_train = []
        y_pred_train = []

        y_val = []
        y_pred_val = []

        y_test = []
        y_pred_test = []

        for data in train_loader:
            inputs, labels = data
            pred = model(inputs)
            y_train.extend(labels.detach().numpy().flatten())
            y_pred_train.extend(pred.detach().numpy().flatten())

        # Iterate through batches
        for data in val_loader:
            inputs, labels = data
            pred = model(inputs)
            y_val.extend(labels.detach().numpy().flatten())
            y_pred_val.extend(pred.detach().numpy().flatten())

        for data in test_loader:
            inputs, labels = data
            pred = model(inputs)
            y_test.extend(labels.detach().numpy().flatten())
            y_pred_test.extend(pred.detach().numpy().flatten())

    y_pred_train = np.array(y_pred_train)
    y_pred_val = np.array(y_pred_val)
    y_pred_test = np.array(y_pred_test)
    y_train = np.array(y_train)
    y_val = np.array(y_val)
    y_test = np.array(y_test)
    train_res = classification_report(y_train, y_pred_train)
    validation_res = classification_report(y_val, y_pred_val)
    test_res = classification_report(y_test, y_pred_test)

    results = pd.DataFrame({"train": train_res, "validate": validation_res, "test": test

    predictions = {
        "train": {"y": y_train, "pred": y_pred_train},

```

```

        "validate": {"y": y_val, "pred": y_pred_val},
        "test": {"y": y_test, "pred": y_pred_test},
    }

    return results, predictions, model

```

## Load the data

```

In [4]: transform = transforms.Compose([transforms.ToTensor(),
                                       transforms.Normalize((0.1307,), (0.3081,))])

train_data = MNIST('./data', train=True, download=True, transform=transform)
test_data = MNIST('./data', train=False, download=True, transform=transform)

```

## Study the MNIST-2D dataset

Length of the dataset

```

In [5]: print("Length of the training set:", len(train_data))
        print("Length of the test set:", len(test_data))
        print()
        print("Shape of \"features\", i.e. images: ", train_data[0][0].shape)

```

```

Length of the training set: 60000
Length of the test set: 10000

```

```

Shape of "features", i.e. images:  torch.Size([1, 28, 28])

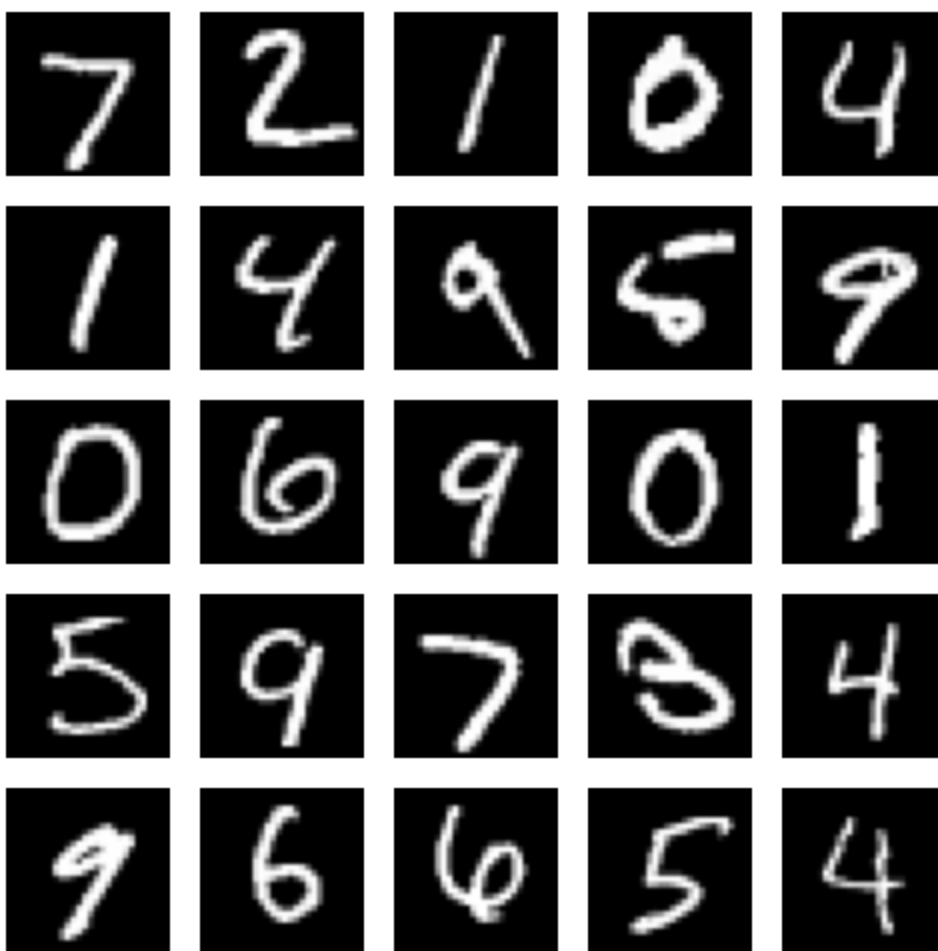
```

## Show some example images

```

In [6]: fig, axs = plt.subplots(5, 5, figsize=(5, 5))
        for i in range(25):
            x, _ = test_data[i]
            ax = axs[i // 5][i % 5]
            ax.imshow(x.view(28, 28), cmap='gray')
            ax.axis('off')
            ax.axis('off')
        plt.tight_layout()
        plt.show()

```



In [7]: *# Distribution of labels in train and test sets*

```
train_labels = [train_data[i][1] for i in range(len(train_data))]
test_labels = [test_data[i][1] for i in range(len(test_data))]

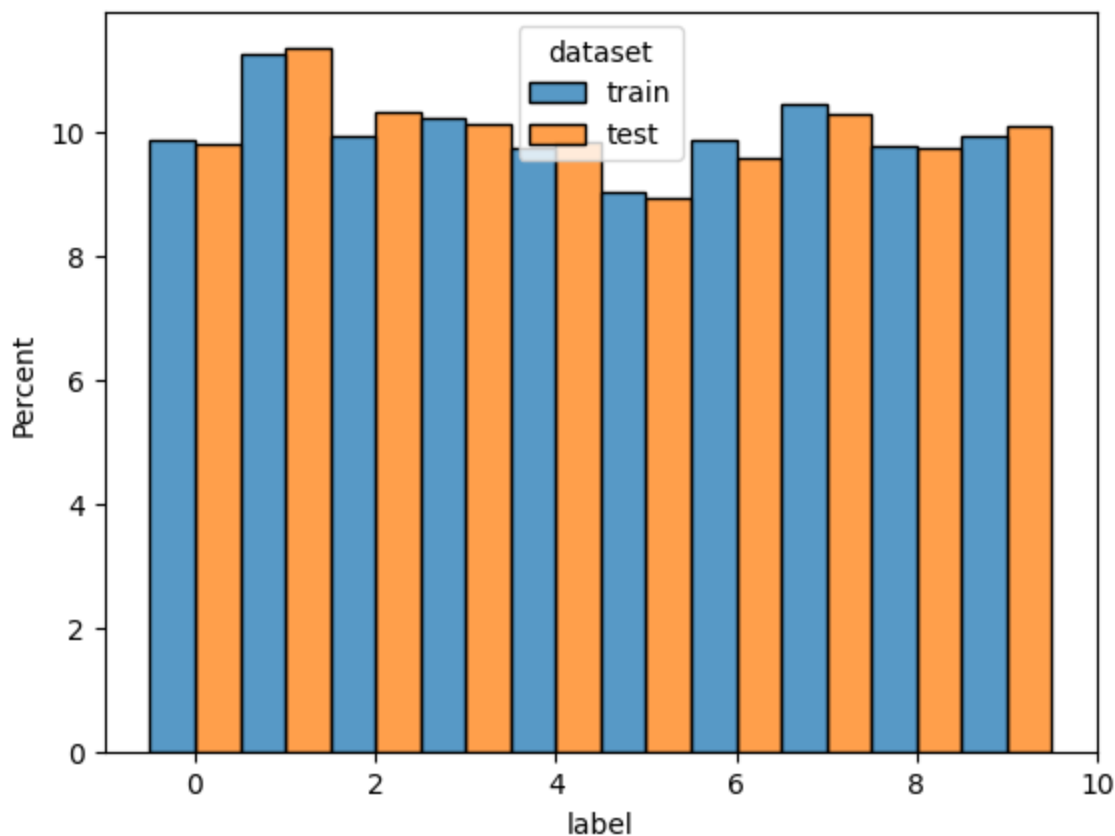
test_features = torch.stack([test_data[i][0] for i in range(len(test_data))])
train_features = torch.stack([train_data[i][0] for i in range(len(train_data))])

train_labels = pd.DataFrame({"label": train_labels, "dataset": "train"})
test_labels = pd.DataFrame({"label": test_labels, "dataset": "test"})
```

In [8]: labels = pd.concat([train\_labels, test\_labels], axis=0)

```
sns.histplot(
    data=labels,
    x="label",
    hue="dataset",
    multiple="dodge",
    stat="percent",
    discrete=True,
    common_norm=False,
)
```

Out[8]: <AxesSubplot: xlabel='label', ylabel='Percent'>



```
In [9]: print("Feature statistics")
print("Min:", train_features.min())
print("Max:", train_features.max())
print("Mean:", train_features.mean())
print("Std:", train_features.std())
```

```
Feature statistics
Min: tensor(-0.4242)
Max: tensor(2.8215)
Mean: tensor(-0.0001)
Std: tensor(1.0000)
```

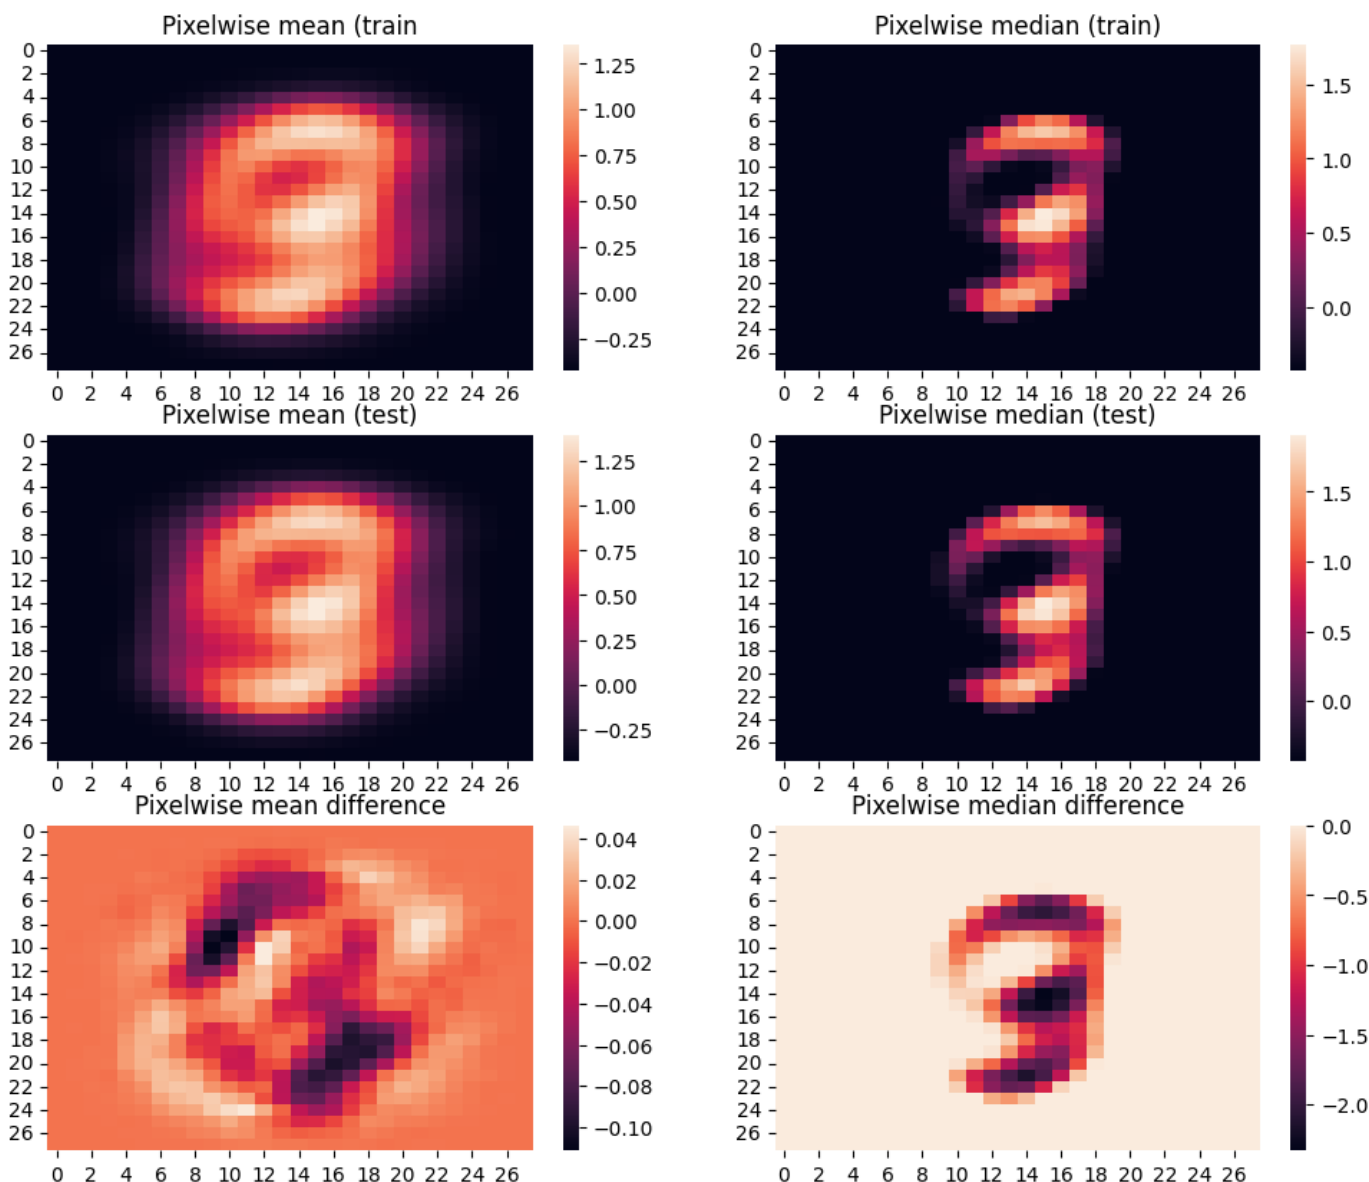
The data is normalized correctly as we see.

## Visualize the pixel distributions in training set

```
In [10]: fig,axs = plt.subplots(3,2, figsize=(12,10))
plt.suptitle("Qualitative feature comparison")
sns.heatmap(train_features.mean(dim=0)[0], ax=axs[0,0])
axs[0,0].set_title("Pixelwise mean (train)")
sns.heatmap(np.median(train_features.detach().numpy(), axis=0)[0], ax=axs[0,1])
axs[0,1].set_title("Pixelwise median (train)")
sns.heatmap(test_features.mean(dim=0)[0], ax=axs[1,0])
axs[1,0].set_title("Pixelwise mean (test)")
sns.heatmap(np.median((test_features).detach().numpy(), axis=0)[0], ax=axs[1,1])
axs[1,1].set_title("Pixelwise median (test)")
sns.heatmap(train_features.mean(dim=0)[0] - test_features.mean(dim=0)[0], ax=axs[2,0])
axs[2,0].set_title("Pixelwise mean difference")
sns.heatmap(np.median(train_features.detach().numpy()) - np.median((test_features).detach().numpy()), ax=axs[2,1])
axs[2,1].set_title("Pixelwise median difference")
```

```
Out[10]: Text(0.5, 1.0, 'Pixelwise median difference')
```





## Feature engineering

For the more traditional models in comparison, the total number of input features  $28 \times 28 = 784$  is significantly high, and PCA is used to reduce the features

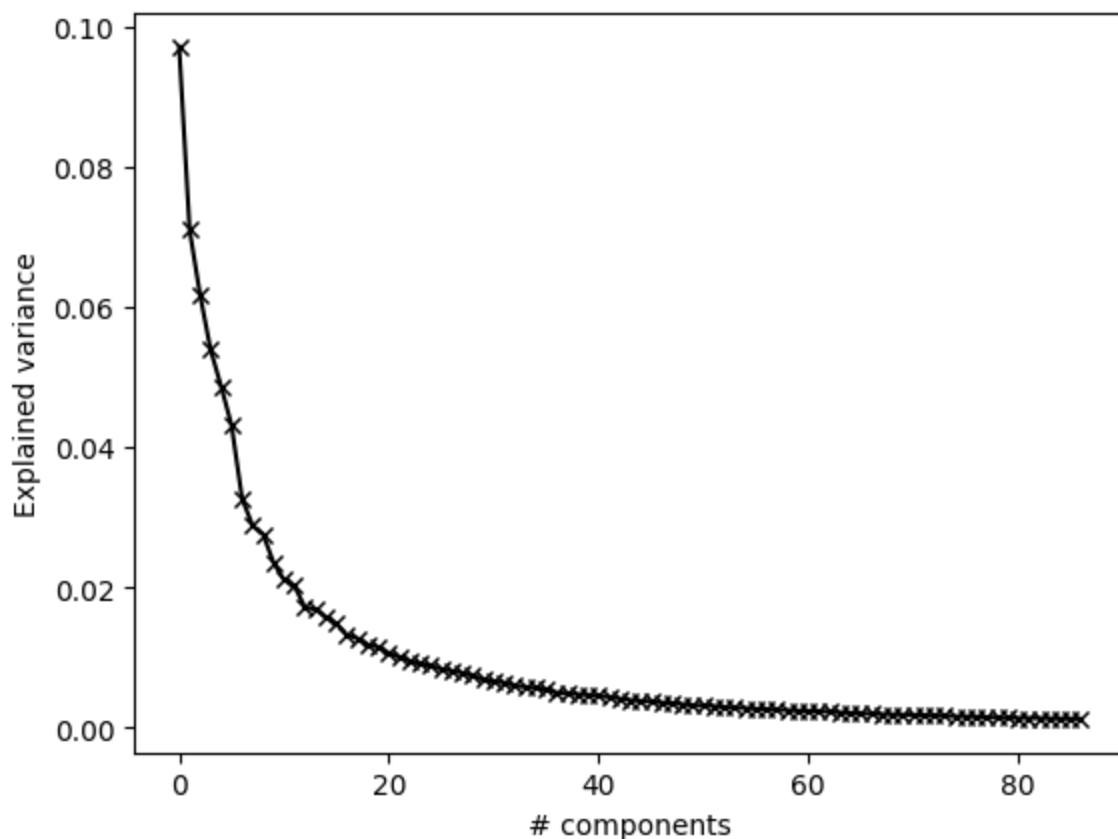
```
In [11]: train_set_flat = train_features.flatten(start_dim=-2).squeeze()
test_set_flat = test_features.flatten(start_dim=-2).squeeze()
train_set_flat.shape, test_set_flat.shape
```

```
Out[11]: (torch.Size([60000, 784]), torch.Size([10000, 784]))
```

## PCA Elbow Curve

```
In [12]: pca = PCA(0.9)
pca.fit(train_set_flat)
plt.plot(pca.explained_variance_ratio_, 'k-x')
plt.xlabel("# components")
plt.ylabel("Explained variance")
```

```
Out[12]: Text(0, 0.5, 'Explained variance')
```



We can see that the curve begins to flatten out after 20 components, let's use that.

## Logistic Regression Model

```
In [13]: model = Pipeline(
    [
        ("pca", PCA(n_components=20)),
        ("reg", LogisticRegression(multi_class='multinomial', fit_intercept=True, max_iter=500))
    ]
)
X_train = train_set_flat
y_train = train_labels["label"]

X_test = test_set_flat
y_test = test_labels["label"]

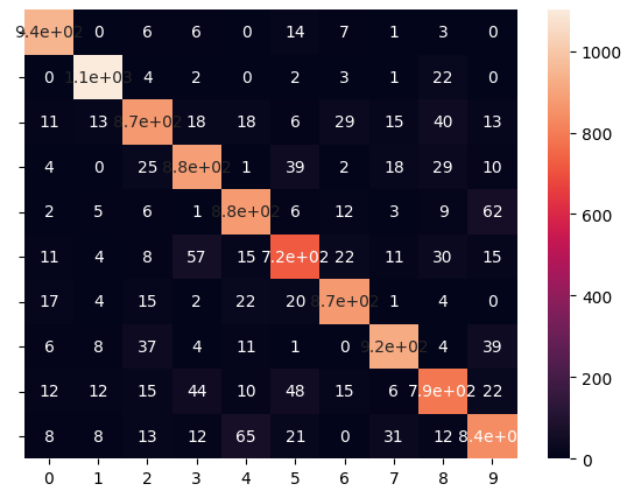
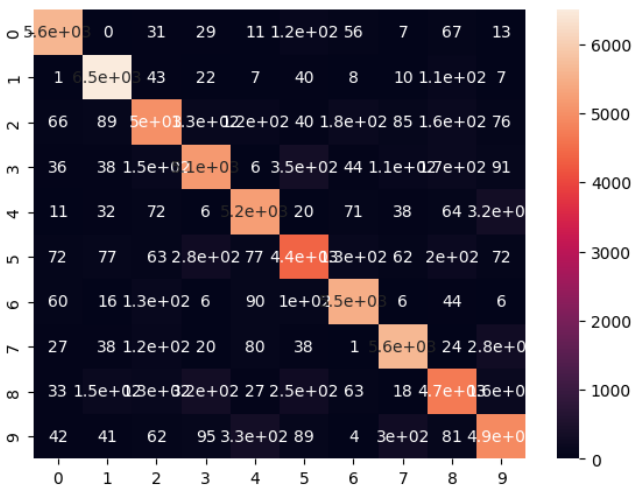
model.fit(X_train, y_train)
y_pred_train = model.predict(X_train)
y_pred_test = model.predict(X_test)

# Additional
print("Train accuracy", accuracy_score(y_train, y_pred_train))
print("Test accuracy", accuracy_score(y_test, y_pred_test))

fig, axs = plt.subplots(1,2, figsize=(15, 5), sharex=True, sharey=True)
sns.heatmap(confusion_matrix(y_train, y_pred_train), annot=True, ax=axs[0])
sns.heatmap(confusion_matrix(y_test, y_pred_test), annot=True, ax=axs[1])

Train accuracy 0.8750666666666667
Test accuracy 0.881
```

```
Out[13]: <AxesSubplot: >
```



## Decision Tree

```
In [14]: def optimize_randomforest(X_train, y_train, random_grid, hparam_max_evals, metric, kfold
# Inspiration from https://towardsdatascience.com/hyperparameter-tuning-the-random-f
# First create the base model to tune
rf = DecisionTreeClassifier()
# Random search of parameters, using k-fold cross validation,
# search across 100 different combinations, and use all available cores
rf_random = RandomizedSearchCV(estimator = rf, param_distributions = random_grid, n_
# Fit the random search model
rf_random.fit(X_train, y_train)
return rf_random
```

Note: As only the training data is used in the randomized hyperparameter search cross validation, it can be considered as using a separate train-val-test split, in which the train and validation sets are used to optimize the model and the test set is used only for testing to avoid overfitting by hyperparameter optimization.

```
In [15]: # Number of features to consider at every split
max_features = np.linspace(0.33, 1.0, 5)
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(5, 10, num = 5)]
max_depth.append(None)
# Minimum number of samples required to split a node
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Create the random grid
random_grid = {
    'max_features': max_features,
    'max_depth': max_depth,
    'min_samples_split': min_samples_split,
    'min_samples_leaf': min_samples_leaf
}

hparam_cv = optimize_randomforest(X_train, y_train, random_grid=random_grid, hparam_max_e
```

Fitting 3 folds for each of 10 candidates, totalling 30 fits

[CV] END max\_depth=5, max\_features=0.8325, min\_samples\_leaf=2, min\_samples\_split=2; total time= 5.3s

[CV] END max\_depth=5, max\_features=0.8325, min\_samples\_leaf=2, min\_samples\_split=2; total time= 5.2s

[CV] END max\_depth=5, max\_features=0.8325, min\_samples\_leaf=2, min\_samples\_split=2; total time= 5.3s

[CV] END max\_depth=7, max\_features=0.665, min\_samples\_leaf=4, min\_samples\_split=10; total time= 5.8s

[CV] END max\_depth=7, max\_features=0.665, min\_samples\_leaf=4, min\_samples\_split=10; total time= 5.5s

[CV] END max\_depth=7, max\_features=0.665, min\_samples\_leaf=4, min\_samples\_split=10; total time= 5.4s

[CV] END max\_depth=6, max\_features=0.8325, min\_samples\_leaf=4, min\_samples\_split=5; total time= 7.1s

[CV] END max\_depth=6, max\_features=0.8325, min\_samples\_leaf=4, min\_samples\_split=5; total time= 7.6s

[CV] END max\_depth=6, max\_features=0.8325, min\_samples\_leaf=4, min\_samples\_split=5; total time= 8.6s

[CV] END max\_depth=10, max\_features=0.4975, min\_samples\_leaf=4, min\_samples\_split=5; total time= 8.4s

[CV] END max\_depth=7, max\_features=1.0, min\_samples\_leaf=1, min\_samples\_split=5; total time= 11.5s

[CV] END max\_depth=7, max\_features=1.0, min\_samples\_leaf=1, min\_samples\_split=5; total time= 11.6s

[CV] END max\_depth=8, max\_features=0.33, min\_samples\_leaf=1, min\_samples\_split=10; total time= 4.3s

[CV] END max\_depth=10, max\_features=0.4975, min\_samples\_leaf=4, min\_samples\_split=5; total time= 8.1s

[CV] END max\_depth=7, max\_features=1.0, min\_samples\_leaf=1, min\_samples\_split=5; total time= 11.2s

[CV] END max\_depth=8, max\_features=0.33, min\_samples\_leaf=1, min\_samples\_split=10; total time= 3.5s

[CV] END max\_depth=8, max\_features=0.33, min\_samples\_leaf=1, min\_samples\_split=10; total time= 3.1s

[CV] END max\_depth=10, max\_features=0.4975, min\_samples\_leaf=4, min\_samples\_split=5; total time= 5.6s

[CV] END max\_depth=6, max\_features=0.33, min\_samples\_leaf=1, min\_samples\_split=2; total time= 2.2s

[CV] END max\_depth=6, max\_features=0.33, min\_samples\_leaf=1, min\_samples\_split=2; total time= 2.8s

[CV] END max\_depth=6, max\_features=0.33, min\_samples\_leaf=1, min\_samples\_split=2; total time= 2.9s

[CV] END max\_depth=10, max\_features=0.8325, min\_samples\_leaf=1, min\_samples\_split=10; total time= 8.6s

[CV] END max\_depth=10, max\_features=0.8325, min\_samples\_leaf=1, min\_samples\_split=10; total time= 8.5s

[CV] END max\_depth=10, max\_features=0.8325, min\_samples\_leaf=1, min\_samples\_split=10; total time= 8.4s

[CV] END max\_depth=8, max\_features=0.665, min\_samples\_leaf=2, min\_samples\_split=10; total time= 5.4s

[CV] END max\_depth=8, max\_features=0.665, min\_samples\_leaf=2, min\_samples\_split=10; total time= 5.1s

[CV] END max\_depth=8, max\_features=0.665, min\_samples\_leaf=2, min\_samples\_split=10; total time= 4.8s

[CV] END max\_depth=None, max\_features=0.665, min\_samples\_leaf=2, min\_samples\_split=5; total time= 9.3s

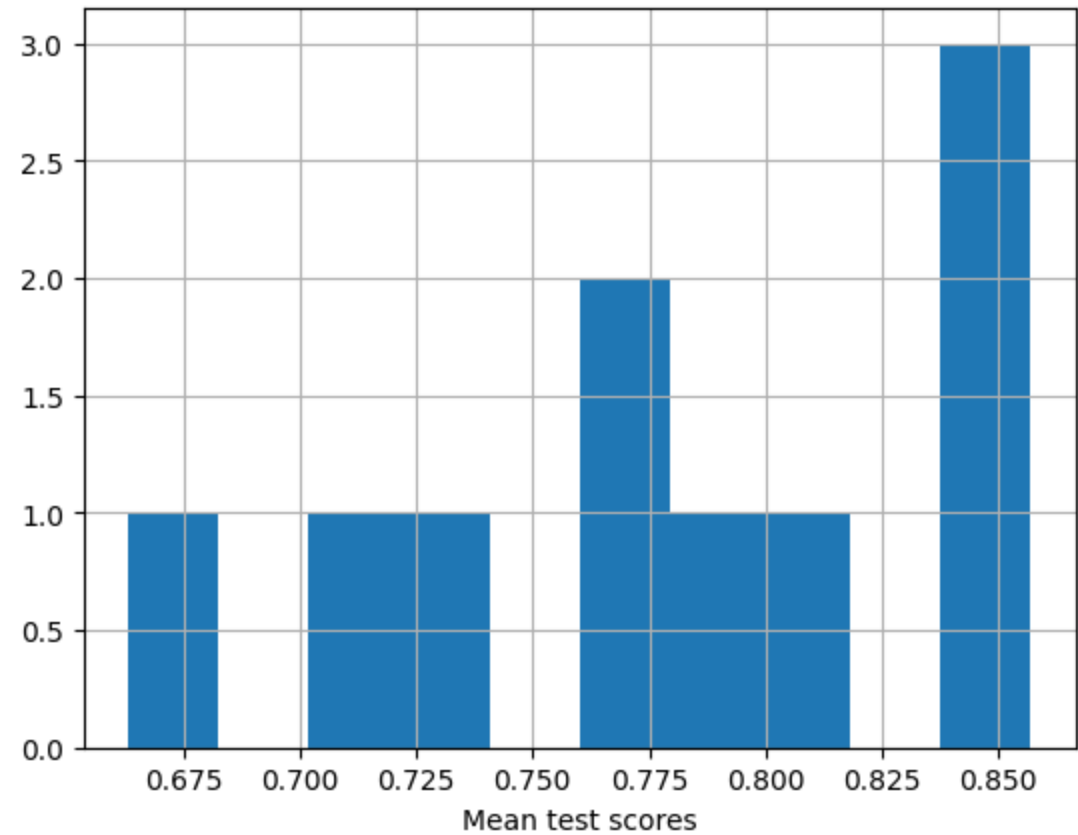
[CV] END max\_depth=None, max\_features=0.665, min\_samples\_leaf=2, min\_samples\_split=5; total time= 9.6s

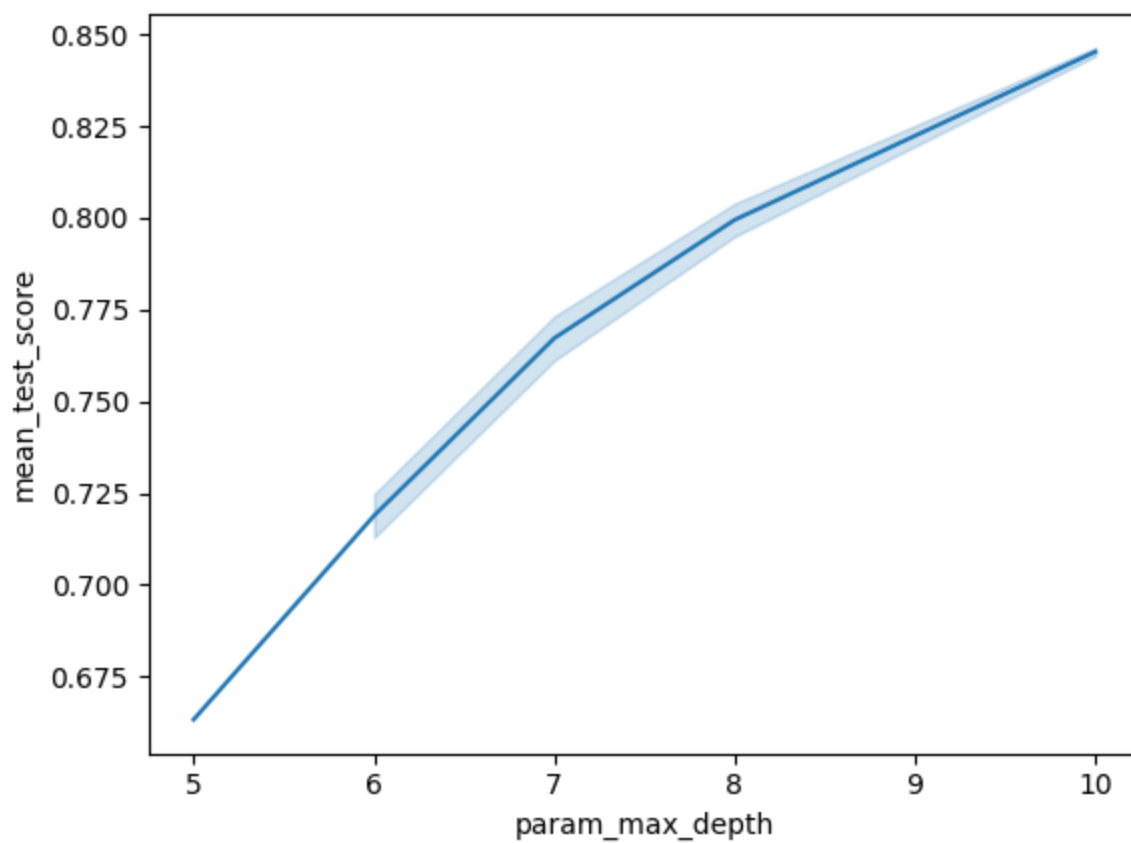
[CV] END max\_depth=None, max\_features=0.665, min\_samples\_leaf=2, min\_samples\_split=5; total time= 8.8s

```
In [16]: hparam_data = pd.DataFrame(hparam_cv.cv_results_)
display(hparam_data[["mean_test_score", "mean_fit_time"]].describe().round(3))
hparam_data["mean_test_score"].hist()
plt.xlabel("Mean test scores")
plt.figure()
sns.lineplot(data=hparam_data, x="param_max_depth", y="mean_test_score", estimator="mean"
```

	mean_test_score	mean_fit_time
count	10.000	10.000
mean	0.778	6.543
std	0.064	2.673
min	0.663	2.578
25%	0.734	5.090
50%	0.784	6.315
75%	0.834	8.201
max	0.857	11.278

```
Out[16]: <AxesSubplot: xlabel='param_max_depth', ylabel='mean_test_score'>
```





We can see that the maximum value of max depth is limiting our performance. However as the parameter also increases fitting time, let us first apply the PCA.

```
In [17]: pca = PCA(n_components=20)
hparam_cv_pca = optimize_randomforest(pca.fit_transform(X_train), y_train, random_grid=r
```

[illegible]

[illegible]



[illegible]

```

al time=    0.6s
[CV] END max_depth=6, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time=    0.5s
[CV] END max_depth=6, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time=    0.5s
[CV] END max_depth=7, max_features=0.4975, min_samples_leaf=2, min_samples_split=10; total time=    0.5s
[CV] END max_depth=7, max_features=0.4975, min_samples_leaf=2, min_samples_split=10; total time=    0.5s
[CV] END max_depth=6, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time=    0.4s
[CV] END max_depth=6, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time=    0.4s
[CV] END max_depth=6, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time=    0.4s

```

```

In [18]: hparam_data_pca = pd.DataFrame(hparam_cv_pca.cv_results_)
display(hparam_data_pca[["mean_test_score", "mean_fit_time"]].describe().round(3))
hparam_data_pca["mean_test_score"].hist()
plt.xlabel("Mean test scores")
plt.figure()
sns.lineplot(data=hparam_data_pca, x="param_max_depth", y="mean_test_score", estimator="

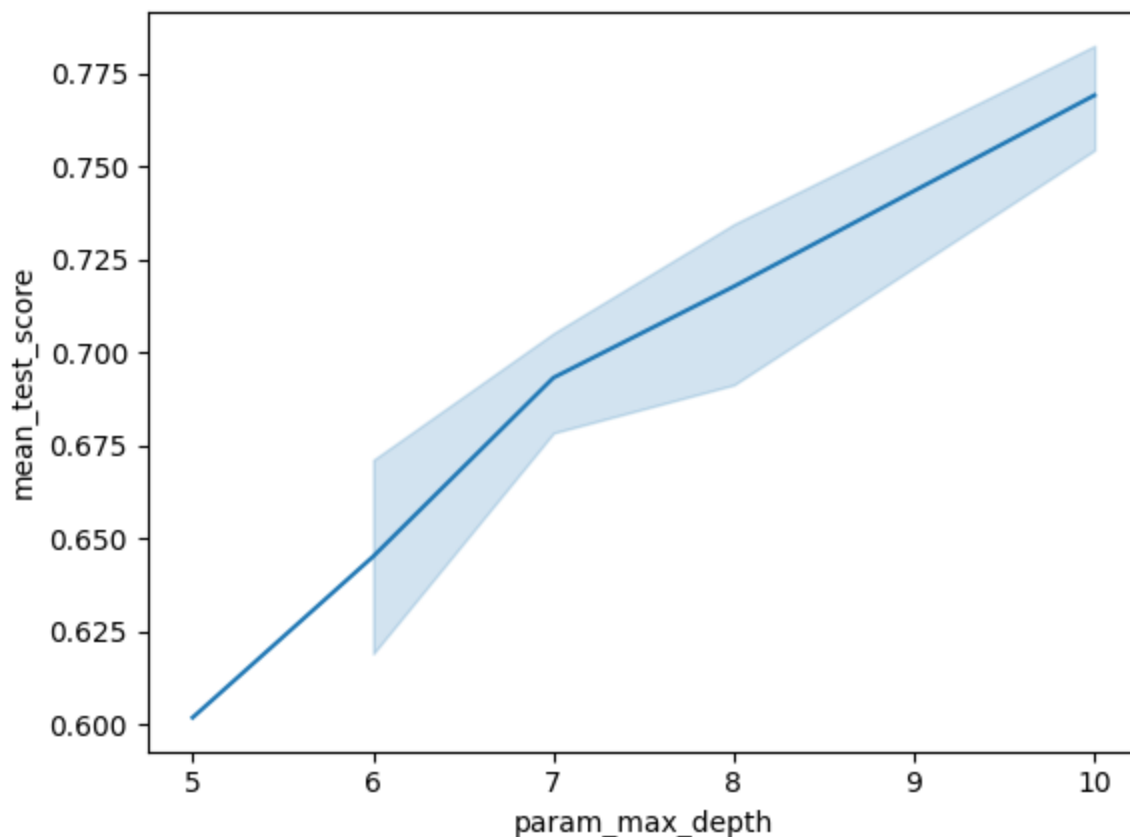
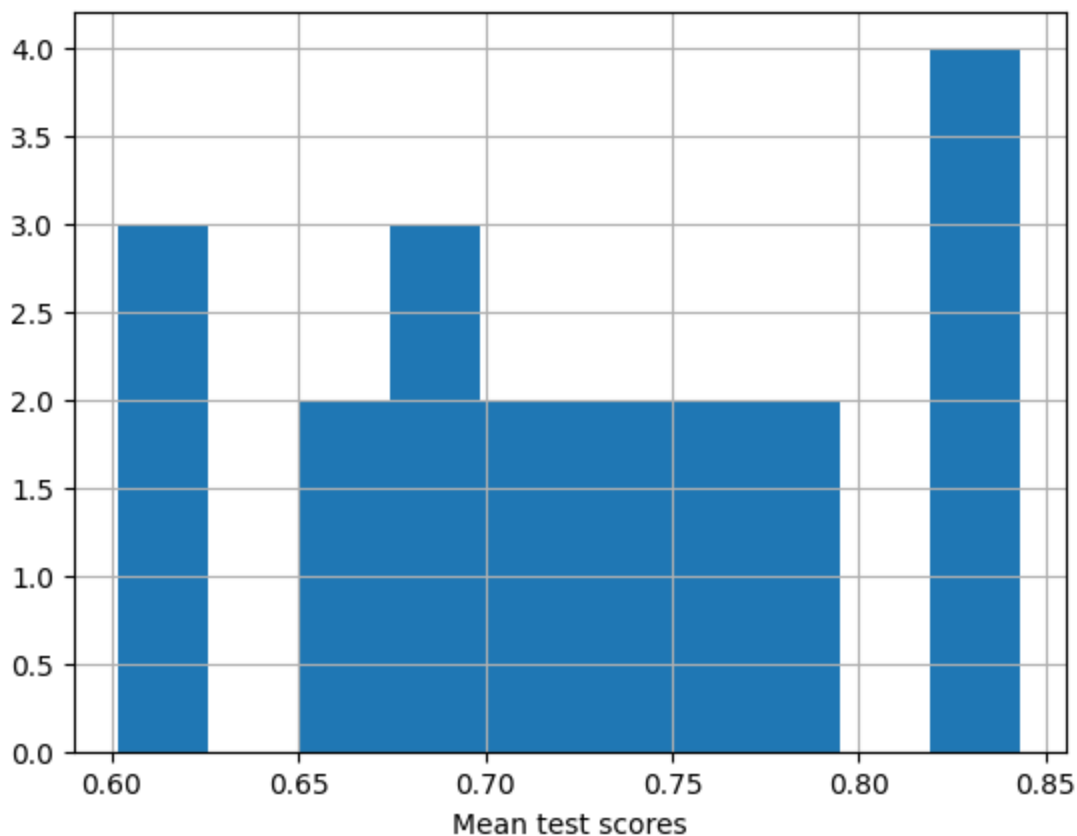
```

	mean_test_score	mean_fit_time
count	20.000	20.000
mean	0.727	0.921
std	0.074	0.553
min	0.602	0.338
25%	0.677	0.634
50%	0.725	0.787
75%	0.778	1.008
max	0.843	2.762

```

Out[18]: <AxesSubplot: xlabel='param_max_depth', ylabel='mean_test_score'>

```



As we can see the performance didn't decrease even though the number of features is lower after the PCA. Let us increase the depth, since the training is significantly faster with the PCA.

```
In [19]: # Number of features to consider at every split
max_features = np.linspace(0.33, 1.0, 5)
# Maximum number of levels in tree
max_depth = [int(x) for x in np.linspace(10, 25, num = 5)]
max_depth.append(None)
# Minimum number of samples required to split a node
```

```
min_samples_split = [2, 5, 10]
# Minimum number of samples required at each leaf node
min_samples_leaf = [1, 2, 4]
# Create the random grid
random_grid = {
    'max_features': max_features,
    'max_depth': max_depth,
    'min_samples_split': min_samples_split,
    'min_samples_leaf': min_samples_leaf
}

pca = PCA(n_components=20)
hparam_cv_pca = optimize_randomforest(pca.fit_transform(X_train), y_train, random_grid=r
```



[illegible]

[illegible]

```

tal time= 1.1s
[CV] END max_depth=17, max_features=0.4975, min_samples_leaf=2, min_samples_split=10; total time= 1.1s
[CV] END max_depth=17, max_features=0.4975, min_samples_leaf=2, min_samples_split=10; total time= 1.1s
[CV] END max_depth=13, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time= 0.9s
[CV] END max_depth=13, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time= 0.9s
[CV] END max_depth=13, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time= 0.8s
[CV] END max_depth=13, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time= 0.7s
[CV] END max_depth=13, max_features=0.4975, min_samples_leaf=4, min_samples_split=2; total time= 0.7s

```

```

In [20]: hparam_data_pca = pd.DataFrame(hparam_cv_pca.cv_results_)
display(hparam_data_pca[["mean_test_score", "mean_fit_time"]].describe().round(3))
hparam_data_pca["mean_test_score"].hist()
plt.xlabel("Mean test scores")
plt.figure()
sns.lineplot(data=hparam_data_pca, x="param_max_depth", y="mean_test_score", estimator="

```

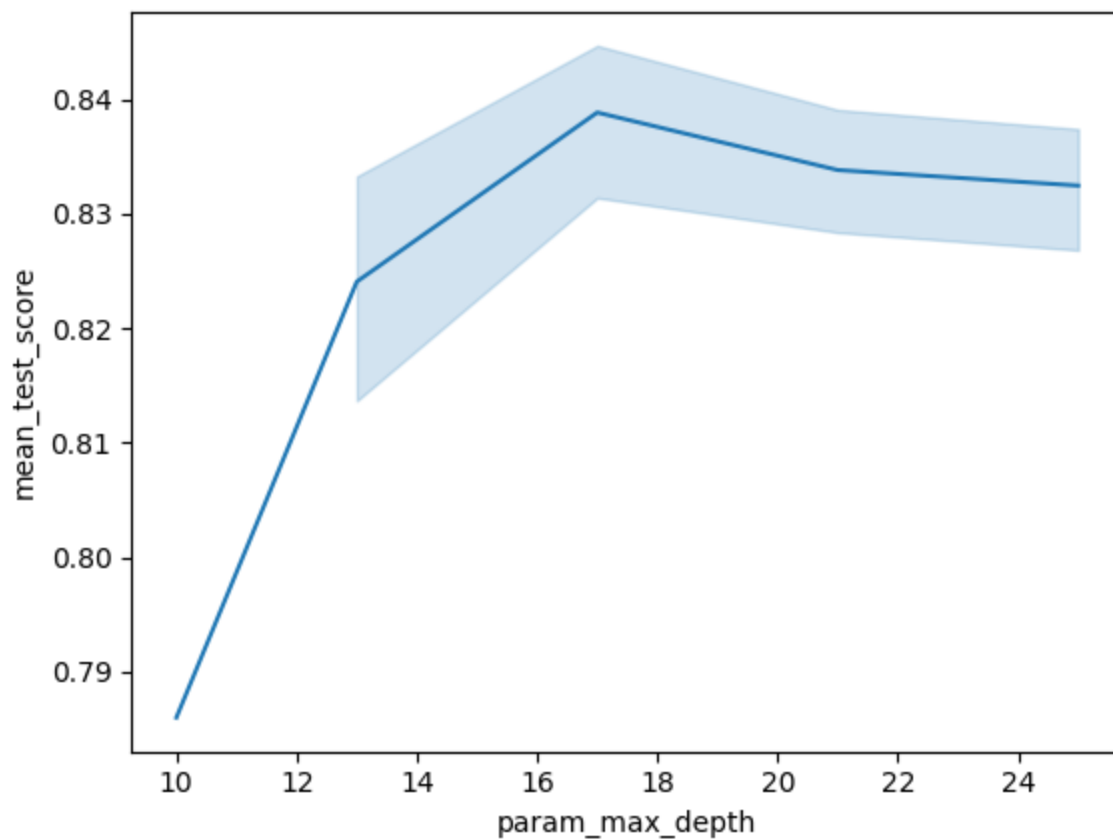
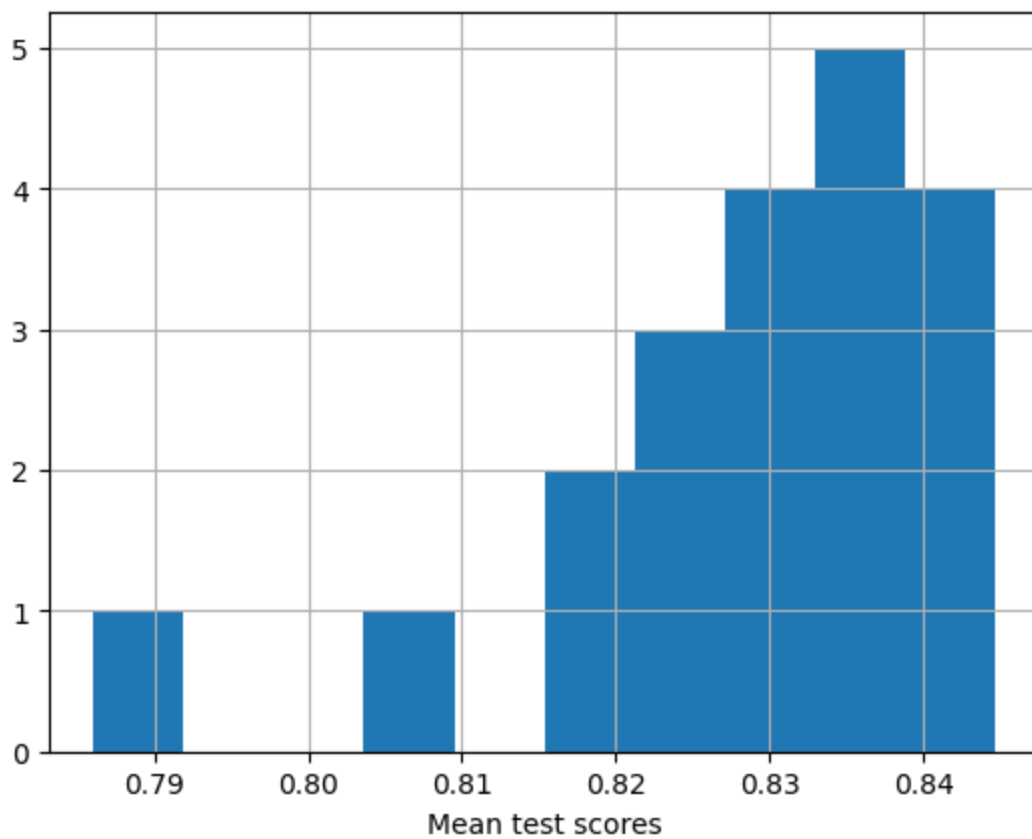
	mean_test_score	mean_fit_time
count	20.000	20.000
mean	0.829	1.554
std	0.014	0.722
min	0.786	0.624
25%	0.823	0.951
50%	0.832	1.361
75%	0.838	2.080
max	0.845	3.304

```

Out[20]: <AxesSubplot: xlabel='param_max_depth', ylabel='mean_test_score'>

```





## Decision Tree Full Training Set test

```
In [21]: model_dt = Pipeline(
    [
        ("pca", PCA(n_components=20)),
        ("clf", DecisionTreeClassifier(**hparam_cv.best_params_))
    ]
)
```

```

model_dt.fit(X_train, y_train)
y_pred_train = model_dt.predict(X_train)
y_pred_test = model_dt.predict(X_test)

# Additional
print("Train accuracy", accuracy_score(y_train, y_pred_train))
print("Test accuracy", accuracy_score(y_test, y_pred_test))

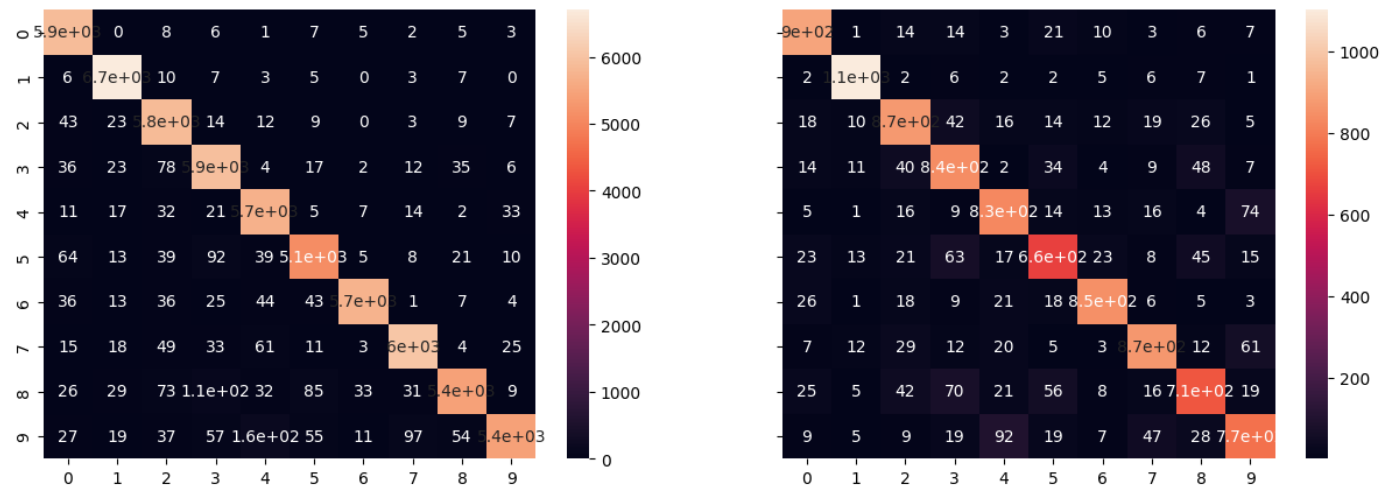
fig, axs = plt.subplots(1,2, figsize=(15, 5), sharex=True, sharey=True)
sns.heatmap(confusion_matrix(y_train, y_pred_train), annot=True, ax=axs[0])
sns.heatmap(confusion_matrix(y_test, y_pred_test), annot=True, ax=axs[1])

```

Train accuracy 0.9629333333333333

Test accuracy 0.8412

Out[21]: <AxesSubplot: >



## KMeans

Let us see if the classes are separable in the latent space

```

In [22]: from sklearn.cluster import KMeans
from sklearn.decomposition import PCA

pca = PCA(n_components=20)
kmeans = KMeans(n_clusters=10) # 1 for each label

pipeline = Pipeline([
    ("pca", pca),
    ("kmeans", kmeans)
])

cluster_labels = pipeline.fit_predict(X_train)

pca_X = pca.fit_transform(X_train)

```

```

In [23]: fig, axs = plt.subplots(1,2, figsize=(15,8))
for i in range(10):
    mask = y_train == i
    axs[0].scatter(pca_X[mask, 0], pca_X[mask,1], label=i, marker=".", alpha=.2)
    axs[0].legend()
for i in range(10):
    mask = cluster_labels == i

```

```

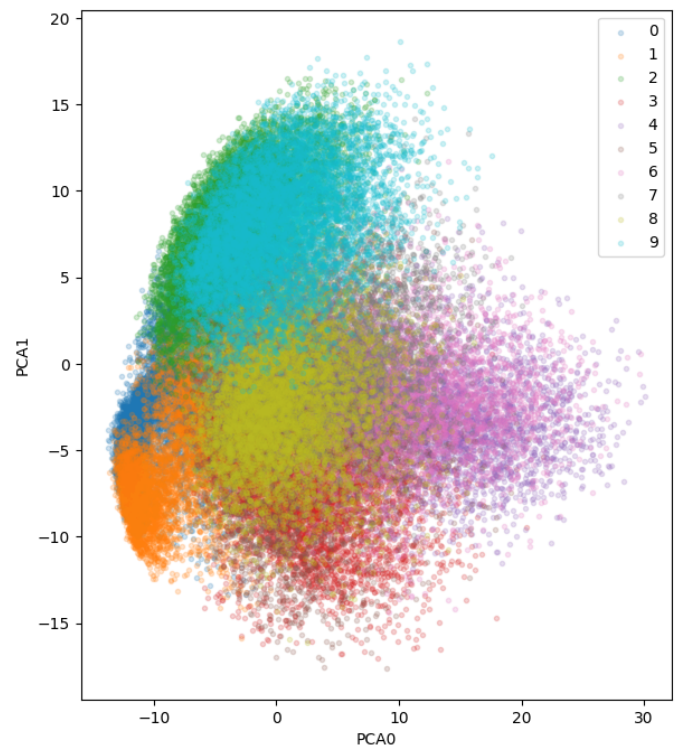
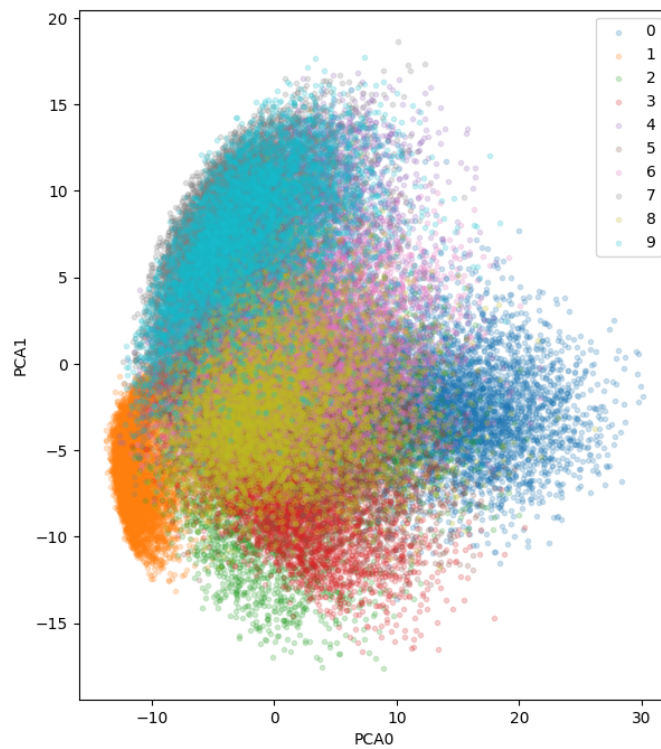
axs[1].scatter(pca_X[mask, 0], pca_X[mask,1], label=i, marker=".", alpha=.2)
axs[1].legend()

```

```

for ax in axs:
    ax.set_xlabel("PCA0")
    ax.set_ylabel("PCA1")

```



```

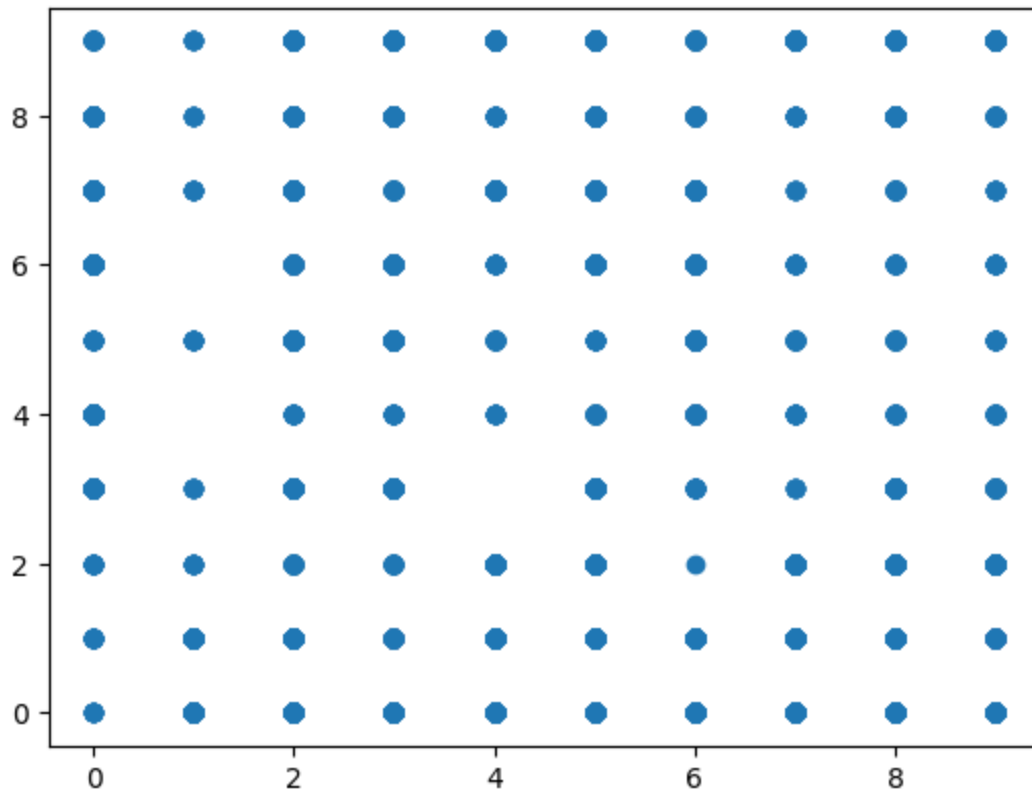
In [24]: plt.scatter(y_train, cluster_labels)

```

```

Out[24]: <matplotlib.collections.PathCollection at 0x7f7bf8dab880>

```



...while similar clusters can be seen, no inference can be made directly on the cluster assignment